


**Standard practice stops at the explanation. This thesis measures whether it can be trusted.**



- 
- Would a second run give the same answer?
  - Do those features actually drive the prediction?
  - Would SHAP name the same three?
  - What does producing this cost?